

Luke Nittmann

luke.nittmann@gmail.com
+1 (313) 500-1244
3038 Isabel Dr, Los Angeles, CA 90026
https://luke-nittmann.vercel.app



EXPERIENCE

Independent Projects (2023 - Present)

Building production-grade applications, focusing on AI-native experiences and real-time collaboration:

Ther <https://ther-app.vercel.app/>

Empathetic AI companion for mental wellness

- Built supportive AI therapist system with personalized chat experience and agentic framework
- Implemented adaptive suggestion engine with contextual intelligence
- Developed secure subscription management with Stripe integration
- Created cross-platform experience with Tauri (iOS/macOS native support)

Tech Stack:

- Frontend: Next.js, React, TypeScript, Tailwind, Framer Motion
- Backend: Prisma, PostgreSQL
- AI/ML: Gemini API, Prompt Engineering
- Payments: Stripe Connect, Payments API
- Cross-platform: Tauri
- State/Utilities: Zustand, Zod, SWR

Cards <https://cards-xyz.vercel.app/>

Personalized job application tracking platform with AI capabilities

- Designed multi-provider agentic system with model abstraction for varied AI responses
- Built embedding-based job-resume matching system for intelligent recommendations
- Generated personalized job application cards from user's links, resume, and web presence
- Implemented context-aware prompt engineering for coherent AI responses
- Created serverless AI processing pipeline with edge function optimization

Tech Stack:

- Frontend: Next.js, React, TypeScript, Tailwind, Framer Motion, shadcn/ui
- Backend: Prisma, NextAuth
- Database: Neon Database, Vercel KV, pgvector
- AI/ML: OpenRouter, Jina, Vercel AI SDK
- State/Utilities: Zustand, SWR, Zod

Voet <https://voet-app.vercel.app/>

Football intelligence platform with AI-powered analysis

- Architected real-time data aggregation system with multi-source orchestration
- Implemented intelligent content analysis with entity recognition using agentic workflows
- Built distributed job queue for resilient scheduled data updates
- Developed relational schema with rich entity relationships

Tech Stack:

- Frontend: Next.js, TypeScript, Tailwind, Framer Motion, Recharts
- Backend: Prisma, Redis, Upstash, tRPC
- Database: PostgreSQL
- AI/ML: Gemini API, Jina, Vector DB, Embeddings
- Data Extraction: Cheerio, Puppeteer, Playwright
- State/Utilities: SWR, Zustand, MSW, Zod

Squish <https://squish-web.vercel.app/>

Semantic social network for content sharing and discovery

- Developed embedding-based content organization system with semantic search
- Built real-time collaborative boards with interactive comments
- Created modular, service-based FastAPI backend with monorepo webapp
- Implemented fluid animations and interactions for enhanced UX

Tech Stack:

- Frontend: Next.js, TypeScript, Tailwind, Motion, Radix UI, react-spring
- Backend: Python, FastAPI, asyncio, pydantic
- Database: PostgreSQL, pgvector, Neon
- Collaboration: Liveblocks, WebSockets
- AI/ML: VertexAI, OpenAI, CLIP, Vector Embeddings
- Infrastructure: Docker, Railway
- Architecture: Turborepo, NX, Module Federation

Loops <https://loops-xyz.vercel.app/>

Stem player for your music with intelligent audio processing

- React / Next.js web app with 4 independent stem track audio loops
- Built AI-powered stem isolation system for vocals, drums, bass, and other
- Implemented client-side audio processing engine with Web Audio API
- Developed efficient audio file streaming and caching system

Tech Stack:

- Frontend: Next.js, React, TypeScript, Tailwind, Framer Motion
- Audio: Tone.js, Web Audio API, AudioWorklet, WebMIDI, FFmpeg.wasm
- Backend: FastAPI, Python
- Storage: Vercel KV, Vercel Blob Storage
- AI/ML: Spleeter
- Infrastructure: Docker, Railway, AWS Lambda
- State: Zustand, Redux Toolkit, Immer

Sine <https://sine-labs.com>

MIDI-based beatmaking app for iOS

- Built native iOS client with comprehensive MIDI engine and patterns
- Developed web-based MIDI/sound pack upload utility
- Implemented AI audio manipulation and semantic search with agentic asset organization
- Created real-time sync with websocket streaming

Tech Stack:

- iOS: Swift, AVFoundation, AudioKit, CoreAudio, MIDIKit, App Intents
- Web: Next.js, TypeScript
- Audio: Pedalboard, FFmpeg, Soundpipe, AudioUnits, VST3
- Database: PostgreSQL, Core Data
- AI/ML: VertexAI, OpenAI, Audio Embeddings
- Cloud: GCP, Cloud Storage, Cloud Run

Demo materials available at luke-nittmann.vercel.app/projects

Titles, Inc.

Senior Software Engineer (Remote) | May 2024 - November 2024

- Led development of cross-platform notification system for web3 image generation platform
- Built semantic image search system with Gemini multimodal embeddings and agentic asset organization
- Architected full-stack features across iOS, web, and backend systems

Technologies:

- Frontend: React, TypeScript, Tailwind CSS, React Three Fiber
- Mobile: Swift, SwiftUI
- Backend: Firebase, Cloud Functions, Firestore
- Cloud: GCP, Cloud Run, Vertex AI
- AI/ML: Gemini API, Multimodal Embeddings, Stable Diffusion, Flux, ComfyUI
- Web3: Ethereum, IPFS, Smart Contracts
- DevOps: CI/CD, Docker

Stems Labs

Senior Software Engineer (Remote) | November 2022 - Present

- Led creation and development of Stem Studio, a mobile audio remixing platform
- Architected full-stack system for professional audio processing with intelligent agents for effects recommendation
- Integrated Python audio models and CoreML for on-device audio processing
- Optimized processing costs through AI-powered effects generation
- Built integration layer supporting major streaming platforms (Spotify, Apple Music)

- Stack:
- Mobile: Swift, AVFoundation, CoreAudio, AudioKit, MIDI
 - Frontend: React, WebGL
 - Backend: Python, Golang, FastAPI, GraphQL, Redis
 - AI/ML: OpenAI, Stable Audio, Demucs, Pedalboard, CoreML, PyTorch
 - Audio: VST3, AudioUnits, FFmpeg, Pedalboard
 - Cloud: AWS, S3, Lambda, EC2, SQS
 - Database: PostgreSQL

Amazon, Inc.

Software Engineer, Address Intelligence (Portland, OR) | June 2019 - December 2021

- Promoted to SDE II within 18 months
- Architected address management system for delivery route optimization
- Improved cross-region performance through stack regionalization
- Built self-service permission system automating internal workflows
- Maintained high service reliability across multiple regions
- Led migration from monolith to microservices architecture

- Technologies:
- Frontend: React Native, TypeScript, Redux, Jest
 - Cloud: AWS, Lambda, DynamoDB, S3, CloudFormation
 - Backend: Java, Python, Node.js, Redis
 - Database: DynamoDB, Aurora PostgreSQL
 - DevOps: CI/CD, Docker, Terraform, CloudWatch
 - Architecture: Microservices, Event-driven, Serverless

AWS Elemental

Software Engineer, AWS MediaConvert (Portland, OR) | January 2018 - June 2019

- Optimized video transcoding engine for improved performance
- Rewrote Quicktime decoder enhancing metadata parsing accuracy
- Developed automated testing framework for validation components
- Resolved customer issues through direct support and code changes

- Technologies:
- Languages: C++, Python, Ruby
 - Cloud: AWS, EC2, S3, CloudFront
 - Media: FFmpeg, H.264/H.265, VP9, AV1
 - Infrastructure: Docker, Jenkins, Terraform
 - Performance: SIMD, Threading, GPU Acceleration
 - Testing: GTest, Pytest, Selenium

EDUCATION

University of Michigan, Ann Arbor

BSc in Computer Science and German Studies (2017)

SKILLS

Programming Languages

- Swift (5+ years)
- Python (7+ years)
- Java
- Ruby
- TypeScript (6+ years)
- C/C++ (8+ years)
- Golang

Frontend

- React
- Next.js
- SwiftUI
- TailwindCSS
- Framer Motion
- Radix UI
- react-spring
- Recharts
- shadcn/ui
- WebGL
- React Three Fiber
- React Native
- Redux
- Motion

Backend

- FastAPI
- Node.js
- Java
- GraphQL
- Prisma ORM
- NextAuth
- Cloud Functions
- pydantic
- asyncio
- tRPC
- WebSockets
- Firebase
- Firestore

Database

- PostgreSQL
- pgvector
- Redis
- Upstash
- DynamoDB
- Vercel Blob Storage
- Core Data
- Aurora PostgreSQL
- Neon DB
- Vercel KV

Audio / Video

- AVFoundation
- FFmpeg
- AudioKit
- CoreAudio
- AudioUnits
- MIDI
- MIDIKit
- Spleeter
- Demucs
- AudioWorklet
- WebMIDI
- H.264/H.265
- VP9
- AV1
- Web Audio API
- Tone.js
- Pedalboard
- VST3
- Python audio models
- Soundpipe
- Stable Audio

AI/ML

- LLMs
- OpenAI
- Gemini
- Claude
- OpenRouter
- PyTorch
- CLIP
- Multimodal AI
- Semantic search
- Agentic systems
- Stable Diffusion
- Flux
- Vector Embeddings
- VertexAI
- Jina
- CoreML
- Context management
- Prompt engineering
- ComfyUI
- Vercel AI SDK
- Multimodal Embeddings

Cloud & Infrastructure

- AWS
- Lambda
- S3
- EC2
- SQS
- CloudFront
- GCP
- Edge Runtime
- Edge Functions
- Firebase
- Railway
- Serverless architecture
- CloudFormation
- Cloud Run
- Cloud Storage
- Vertex AI
- Vercel
- Distributed systems
- Edge computing
- CloudWatch
- Microservices
- Event-driven

Data & Testing

- Jest
- GTest
- Pytest
- Selenium
- MSW
- Cheerio
- Puppeteer
- Playwright

Tools & Practices

- Git
- Vim
- Docker
- CI/CD
- Monitoring
- Tauri
- Liveblocks
- Module Federation
- NX
- SIMD
- Turborepo
- Monorepo
- Terraform
- Jenkins
- Threading
- GPU Acceleration
- App Intents

State Management

- Zustand
- Redux Toolkit
- Immer
- SWR
- Zod